

Lecture Notes 3: Sampling, Simulation, Low-Level Extensions

Sampling from Distributions

```
In [1]: import numpy, numpy.random
```

Sampling from a uniform distribution between 0 and 1

```
In [2]: print numpy.random.uniform(0,1,[10])
```

```
[ 0.72611138  0.24067817  0.71938139  0.45283454  0.70368496  0.94023131
  0.77130642  0.04720969  0.29922044  0.66994318]
```

Sampling from normal distribution of mean 10 and standard deviation 0.01

```
In [3]: print numpy.random.normal(10,0.01,[10])
```

```
[ 10.00598152  10.00735394  10.00212889   9.9986185   10.00672975
   9.98147864  10.00094005   9.995714   10.00350466  10.00550996]
```

Demo 1: Testing A Statistical Paradox (the James-Stein Paradox)

When estimating the mean of a data distribution, moving the estimator from the empirical mean, and closer to some arbitrary point (e.g. the origin) makes the estimator more accurate. Let's verify it:

```
In [6]: # Testing many times:
        errstd = []
        errstd2 = []

        n = 10
        d = 50

        for i in range(100):

            # sample from a distribution of mean vector 1 and standard deviation 1
            m = 1.0

            X = numpy.random.normal(m,1,[n,d])

            # empirical mean
            m_emp = X.mean(axis=0)

            # some coefficient
            c = (1-(d-2)*1.0/n/((m_emp)**2)).sum()

            # james-stein estimator
            m_js = c*m_emp

            # the error between the true mean and the standard estimator
            errstd += [(m - m_emp)**2].sum()


            # the error between the true mean and the said better estimator
            errstd2 += [(m - m_js)**2].sum()
```

```
In [7]: import matplotlib
        from matplotlib import pyplot as plt
        %matplotlib inline
        from IPython.display import set_matplotlib_formats
        set_matplotlib_formats('pdf','png')
        plt.rcParams['savefig.dpi'] = 90
```

```
In [9]: plt.scatter(errstd,errstd2)

        l,h = min(errstd+errstd2),max(errstd+errstd2)
        plt.plot([l,h],[l,h],color='red')
        plt.xlabel('error for standard mean estimator')
        plt.ylabel('error for James-Stein mean estimator')
        plt.axis([l,h,l,h])
```

```
Out[9]: [2.4355343709629809,
        7.4646057017558283,
        2.4355343709629809,
        7.4646057017558283]
```



Observation: most points are below the curve (i.e. James-Stein estimator is better).

Making discrete choices

Let us suppose we have 7 fruits to choose from:



Figure 1:

```
In [10]: fruits = ['watermelon','apple','grape','grapefruit','lemon','banana','cherry']
```

We use the function `random.choice` to randomly choose from that list

```
In [11]: import random
         print random.choice(fruits)
```

grapefruit

```
In [12]: print random.choice(fruits)
         print random.choice(fruits)
         print random.choice(fruits)
```

banana

lemon

apple

The function choice of the module `numpy.random` provides some more functionalities such as choosing repeatedly

```
In [11]: print numpy.random.choice(fruits,[20])
```

```
['apple' 'lemon' 'apple' 'grape' 'banana' 'grapefruit' 'cherry' 'lemon'
 'watermelon' 'grapefruit' 'banana' 'watermelon' 'grapefruit' 'cherry'
 'cherry' 'banana' 'cherry' 'banana' 'lemon' 'banana']
```

or specify a certain distribution of fruits to choose from

```
In [12]: p = [0.05,0.70,0.05,0.05,0.05,0.05,0.05]
         print numpy.random.choice(fruits,[20],p=p)
```

```
['grapefruit' 'apple' 'apple' 'apple' 'apple' 'apple' 'apple' 'apple'
 'banana' 'apple' 'watermelon' 'cherry' 'apple' 'apple' 'apple' 'apple'
 'watermelon' 'apple' 'cherry' 'apple']
```

Another way to make discrete choices




Figure 2:

```
In [13]: # Define fruits probabilities
         p = [0.05,0.70,0.05,0.05,0.05,0.05]

         # Cumulate them
```

```

l = numpy.cumsum([0]+p[:-1]) # lower-bounds
h = numpy.cumsum(p)           # upper-bounds

# Draw a number between 0 and 1
u = numpy.random.uniform(0,1)

# Find which basket it belongs to
s = (u>l)*(u<h)

# retrieve the label
fruits[numpy.argmax(s)]

```

Out[13]: 'apple'

The code can be parallelized to choose multiple fruits at the same time

```

In [14]: # generate many numbers between 0 and 1
u = numpy.random.uniform(0,1,[20])

# find the basket to which they belong
na = numpy.newaxis
s = (u[:,na]>l[na,:])*(u[:,na]<h[na,:])

print([fruits[i] for i in numpy.argmax(s,axis=1)])

```

['watermelon', 'apple', 'apple', 'apple', 'apple', 'apple', 'apple', 'apple', 'grapefruit', 'apple', 'banana', '']

More importantly, the code can be parallelized to choose from multiple distributions at the same time. Let us create a matrix of probability distributions.

```

In [15]: P = [
    [0.0,0.05,0.05,0.05,0.05,0.05,0.05,0.70],
    [0.0,0.05,0.05,0.05,0.05,0.05,0.70,0.05],
    [0.0,0.05,0.05,0.70,0.05,0.05,0.05,0.05],
]
# (note that we have added a zero-probability state
# at the beginning for ease of implementation)

```

We stack them, and compute the bounds L and H

```

In [16]: C = numpy.cumsum(P,axis=1) # lower-bounds

L,H = C[:,-1],C[:,1:]

```

Draw 8 fruits from each distribution, and test the bounds

```

In [19]: R = numpy.random.uniform(0,1,[3,8])

S = (R[:,na] > L[:,na,:])*(R[:,na] < H[:,na,:])

In [20]: for s in S:
    print([fruits[i] for i in numpy.argmax(s,axis=1)])

```

['lemon', 'cherry', 'cherry', 'cherry', 'cherry', 'apple', 'cherry', 'cherry']
 ['banana', 'banana', 'banana', 'banana', 'grapefruit', 'cherry', 'grapefruit', 'banana']
 ['grape', 'grape', 'grape', 'grape', 'grape', 'grape', 'lemon', 'grape']




Figure 3:

Demo 2: Monte Carlo Simulation of a Markov Chain

A Markov chain transits between a set of states, where the transition between pairs of states is associated with a fixed probability. The set of probabilities can be stored in a transition matrix.

```
In [22]: # Transition matrix
T = numpy.array([
    [0.9,0.1,0.0], # transiting from state 1 to state 1,2,3
    [0.0,0.9,0.1], # transiting from state 2 to state 1,2,3
    [1.0,0.0,0.0], # transiting from state 3 to state 1,2,3
])
```

Thanks to our parallel implementation, we can choose from multiple probability distributions, and thus, simulate multiple markov chains in parallel, each of them being in a different state.

```
In [24]: # Add empty state to transition matrix
P = numpy.pad(T,1,mode='constant')[1:-1,:-1]
print(P)
```

```
[[ 0.  0.9  0.1  0. ]
 [ 0.  0.  0.9  0.1]
 [ 0.  1.  0.  0. ]]
```

```
In [25]: # Implementing transition for each particle
def mcstep(X):
    Xp = numpy.dot(X,P)

    Xc = numpy.cumsum(Xp,axis=1)

    L,H = Xc[:,-1],Xc[:,1:]

    R = numpy.random.uniform(0,1,[len(Xp),1])

    return (R > L)*(R < H)*1.0
```

```
In [26]: # Initialize all particles to state 1
A = numpy.outer(numpy.ones([30]),[1.0,0,0])

print 'initial distribution of particles states: [% .3f % .3f % .3f]'%tuple(A.mean(axis=0))
```

```

for i in range(20):
    print '(iter %2d)'%i, numpy.argmax(A, axis=1)
    A = mcstep(A)

# Print distribution of sampled states
print 'final distribution of particles states: [%.3f %.3f %.3f]'%tuple(A.mean(axis=0))

initial distribution of particles states: [1.000 0.000 0.000]
(iter 0) [0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
(iter 1) [0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
(iter 2) [0 0 1 1 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0]
(iter 3) [0 1 2 1 0 1 0 0 0 0 1 0 0 1 1 0 0 0 0 0 1 0 1 0 0 1 0]
(iter 4) [0 1 0 1 0 1 0 0 0 0 1 0 0 2 1 0 0 0 0 0 1 0 1 0 1 1 0]
(iter 5) [0 1 0 2 1 1 0 0 0 0 1 0 0 0 1 0 0 0 0 0 1 0 1 0 1 1 0]
(iter 6) [0 1 0 0 1 2 0 0 1 0 2 0 0 0 1 0 1 0 0 0 1 0 1 0 1 1 0]
(iter 7) [0 1 0 0 1 0 0 0 2 0 0 0 0 0 1 0 1 0 0 1 1 0 2 0 1 1 0]
(iter 8) [0 1 0 0 1 0 0 0 0 0 0 0 0 0 1 0 1 0 0 1 1 0 0 1 2 1 0]
(iter 9) [0 1 0 0 1 0 0 0 0 1 0 0 0 0 1 0 1 0 0 1 1 0 0 1 0 1 0]
(iter 10) [0 1 0 1 2 0 0 0 0 1 0 0 0 0 1 0 1 0 0 2 1 0 0 1 0 1 0]
(iter 11) [0 1 0 1 0 1 0 0 0 1 0 0 0 0 1 0 2 0 0 0 1 0 0 2 0 1 0]
(iter 12) [0 1 0 1 0 1 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 0 0 0 1 0]
(iter 13) [0 1 0 1 1 1 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 0 0 1 0]
(iter 14) [0 1 0 1 1 2 0 0 0 1 0 0 1 0 1 0 0 0 1 0 2 0 1 0 0 1 0]
(iter 15) [0 1 0 1 1 0 0 0 0 1 0 0 2 0 1 0 0 1 2 0 0 0 1 0 1 1 0]
(iter 16) [0 1 0 1 1 1 0 0 0 1 0 0 0 0 1 0 0 1 0 0 0 0 1 0 1 1 0]
(iter 17) [0 1 0 1 1 2 0 0 0 1 0 0 0 0 1 0 0 1 0 0 0 0 1 0 1 2 0]
(iter 18) [0 1 0 2 1 0 0 0 0 1 0 0 0 0 1 0 0 1 0 0 1 0 1 0 1 0 0]
(iter 19) [0 1 0 0 1 0 0 0 0 1 0 0 0 0 1 0 0 1 0 0 1 0 1 0 1 0 0]
final distribution of particles states: [0.633 0.367 0.000]

```

Low-level extensions

Idea: Write highly optimized functions directly in C/C++, Fortran, or Cuda, and make them accessible to the Python user. Numpy is also based on this principle: making very efficient functions accessible in Python in a user-friendly manner.

Example: F2PY (Fortran to Python)

Main steps:

- Install gfortran (fortran compiler) and f2py (binding between Fortran and Python)
- Create file (e.g. convolution.f90) containing some optimized code for convolutions.
- Compile convolution.f90 into a Python module using f2py
- Import module convolution directly from Python.

Applying batch convolutions

We would like to use the following Fortran code that computes a batch of convolutions in the same way as the convolutional layer of a neural network.

We run the command

```
f2py -c convolution.f90 -m convolution
```

that compiles the code into a module file convolution.so, that we can now access in Python.

```
In [27]: import convolution
```

```
convolution,convolution.conv
```

```
Out[27]: (<module 'convolution' from 'convolution.so'>, <fortran object>)
```

```

! -*- f90 -*-
subroutine conv(x,wx,hx,y,w,ww,hw,nx,ny)
integer wx,hx,nx,wy,hy,ny,ww,hw
real x(wx,hx,nx),y(wx-ww+1,hx-hw+1,ny),w(ww,hw,nx,ny)
!f2py intent(out) :: y
y = 0
wy = wx-ww+1
hy = hx-hw+1
do i=1,nx
do m=1,hw
do k=1,ww
forall(l=1:wy,j=1:ny,n=1:hy) y(l,n,j)=y(l,n,j)+x(l+k-1,n+m-1,i)*w(k,m,i,j)
end do
end do
end do
end

```

Figure 4: file: convolution.f90

Demo 3: Filtering images to detect interesting features

```

In [33]: import numpy,cv2
import matplotlib
from matplotlib import pyplot as plt
%matplotlib inline

In [34]: X = numpy.array(cv2.imread('tea.jpg'))[:, :, ::-1]*1
plt.imshow(X)

Out[34]: <matplotlib.image.AxesImage at 0x7f9d76c85810>

```



```

In [35]: # filter size x
# filter size y
# filter size z
# filter size

W = numpy.zeros([4,4,3,5])

# horizontal edge detector
W[:, :, :, 0] = 1.0
W[:, :, :, 1] = -1.0

```

```

# horizontal edge detector in reverse direction
W[:,2:,:,:1] = -1.0
W[2:,:,:1] = 1.0

# vertical edge detector
W[:,2:,:,2] = 1.0
W[:,2:,:,2] = -1.0

# vertical edge detector in reverse direction
W[:,2:,:,3] = -1.0
W[:,2:,:,3] = 1.0

# more red than yellow detector
W[:,:,0,4] = 1.0
W[:,:,1,4] = -1.0

```

```


In [36]: Z = convolution.conv(X,W)
        Y = numpy.maximum(0,Z)

```

```

In [37]: f = plt.figure(figsize=(10,7))
        for i in range(4):
            p = f.add_subplot(2,2,i+1)
            p.imshow(Y[:,:,:i])

```



```

In [38]: plt.imshow(Y[:, :, 4])

```

```

Out[38]: <matplotlib.image.AxesImage at 0x7f9d8e9ef690>

```